

Multiple Linear Regression (MLR) on Ethereum Prices Dataset

Gabriela Howell

Master of Science Data Analytics, Western Governors University

Data Analytics Graduate Capstone

November 24, 2024

A. Research Question

The research question investigates whether a predictive multiple linear regression (MLR) model can be built for Ethereum's price, utilizing historical price and volume data. The study explores the hypothesis that historical Ethereum data significantly influences its current price, enabling a predictive accuracy of 70% or higher.

- Null hypothesis – There is no significant relationship between historical Ethereum data and current Ethereum prices, with a prediction accuracy lower than 70%.
- Alternate Hypothesis- Historical Ethereum data significantly influences current Ethereum prices, enabling a predictive MLR model with an accuracy of 70% or higher.

This question arises in financial analytics, where cryptocurrency price fluctuations are incredibly volatile. By focusing on Ethereum, the research seeks to demonstrate how traditional financial modeling techniques can provide actionable insights into a highly dynamic market. The study leverages publicly available historical data from Coinbase, offering a foundation for applying MLR to predict short-term price changes and identify key market drivers. McNally, Roche, and Caton (n.d.) argue that historical data, especially in volatile markets, provides crucial insights, supporting the idea that past prices and trading volumes play a significant role in price prediction. If historical data is found to be non-influential, it may suggest that random market forces dominate Ethereum's price fluctuations. The study seeks to advance predictive methodologies in cryptocurrency markets through this analysis.

B. Data Collection

The data for this project will be sourced from Coinbase's public API, which is publicly available on Kaggle [here](#). This dataset includes historical Ethereum pricing data, such as open, close, high, low prices, and volume, and contains over 4.1 million rows, and 6 columns. Which provides an extensive coverage for robust modeling (Bukhari, 2024). This data is freely available and adheres to academic and non-commercial use policies, making it highly accessible. However, the dataset lacks sentiment indicators and macroeconomic variables, which could impact prediction accuracy. Ensuring the data was properly formatted and free of missing values was a challenge, but the dataset's high quality (0.00% sparsity) minimized these issues. Khaniki and Manthouri (2024) highlight preprocessing steps critical to ensuring that machine learning algorithms perform optimally in volatile market conditions. By focusing solely on historical prices and trading volume, the study adopted a streamlined approach suitable for predictive modeling while acknowledging the limitations of excluding broader market factors.

C. Data Extraction and Preparation

Pre-processing steps will include handling any missing values, normalizing the data, and ensuring it meets the assumptions necessary for MLR (Khaniki & Manthouri, 2024). According to Real Python, tools like Python's pandas and NumPy are crucial for effectively cleaning data, addressing missing values, fixing inconsistent formats, and removing outliers (Agarwal, 2024). The dataset, gathered from a publicly available source on Kaggle, is in CSV format. Using

Python's Pandas library, the data was extracted and loaded into a pandas.

```
# Load the Ethereum Dataset
Ethereum = pd.read_csv('ETHUSD_1m_Coinbase.csv')

# Check the first few rows to ensure it loaded correctly
print(Ethereum.head())
```

	Open time	Low	High	Open	Close	Volume
0	2016-09-29 00:00:00	13.30	13.31	13.31	13.30	5.15493
1	2016-09-29 00:02:00	13.30	13.30	13.30	13.30	5.00000
2	2016-09-29 00:03:00	13.33	13.33	13.33	13.33	1.53755
3	2016-09-29 00:05:00	13.30	13.30	13.30	13.30	4.70303
4	2016-09-29 00:06:00	13.33	13.37	13.33	13.37	100.00000

The primary data-preparation steps for this dataset included:

1. **Handling Missing Values:** Although the dataset has a very low proportion of missing values (0.00%), it's essential to verify and handle any missing data before analysis. This step ensures the dataset is complete and ready for modeling. Since there were no missing values, no further action was required.
2. **Checking for Duplicates:** After verifying missing values, the dataset was checked for duplicates, and none were found.

```
# Check for duplicate rows
duplicates = Ethereum[Ethereum.duplicated()]

# Find rows with missing values
rows_with_missing = Ethereum[Ethereum.isnull().any(axis=1)]

# Print summary of duplicates and missing values
print(f"Number of duplicate rows: {len(duplicates)}")
print(f"Number of rows with missing values: {len(rows_with_missing)}")

Number of duplicate rows: 0
Number of rows with missing values: 0
```

3. **Outlier Removal:** Outliers can significantly impact model performance, particularly in financial datasets. Boxplots were utilized to identify and remove outliers. The primary column with outliers was 'Volume'. Analysis revealed a substantial gap between 1,500 and over 120,000, with 400,982 outliers identified. Consequently, any values above 1,500 were removed.

```
# Display the count of outliers
print(f"Number of outliers:")
print(count_outliers)

Number of outliers:
Low          0
High         0
Open         0
Close        0
Volume      400982
dtype: int64

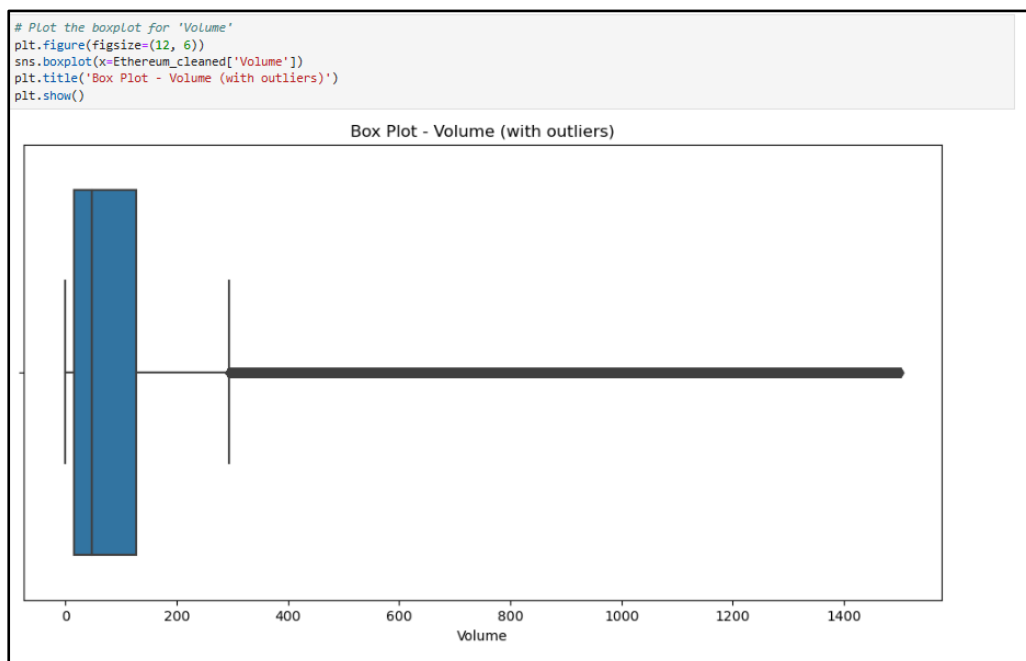
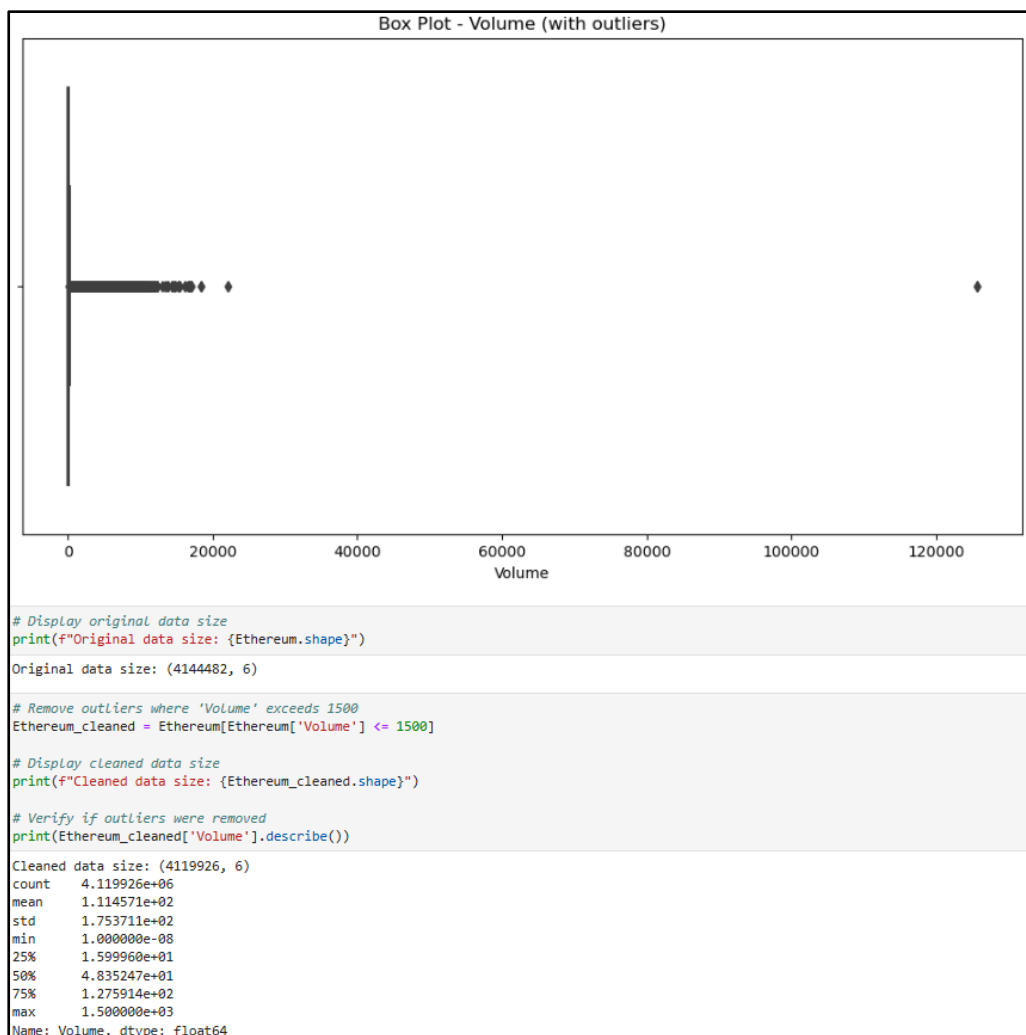
# Calculate descriptive statistics for 'Volume'
volume_stats = Ethereum['Volume'].describe()

# Identify outliers in the 'Volume' column
q1_volume = volume_stats['25%']
q3_volume = volume_stats['75%']
iqr_volume = q3_volume - q1_volume
lower_whisker_volume = q1_volume - 1.5 * iqr_volume
upper_whisker_volume = q3_volume + 1.5 * iqr_volume
outliers_volume = (Ethereum['Volume'] < lower_whisker_volume) | (Ethereum['Volume'] > upper_whisker_volume)

print(f"Column: Volume")
print(f"Lower Whisker: {lower_whisker_volume}")
print(f"Upper Whisker: {upper_whisker_volume}")
print(f"Count of Outliers: {outliers_volume.sum()}")

# Plot the boxplot for 'Volume'
plt.figure(figsize=(12, 6))
sns.boxplot(x=Ethereum['Volume'])
plt.title('Box Plot - Volume (with outliers)')
plt.show()

Column: Volume
Lower Whisker: -154.7764581
Upper Whisker: 300.97249462
Count of Outliers: 400982
```



The dataset exhibits an exceptionally low proportion of missing values, with a sparsity of 0.00%, highlighting its high-quality nature. This makes it ideal for conducting reliable analysis and creating accurate predictive models for Ethereum pricing. Since all variables are already quantitative, no further data conversion is required. After the data was cleaned, I saved the data 'cleaned_data.csv' which I will include with the file.

```
# Calculate sparsity
total_elements = Ethereum.size
missing_elements = Ethereum.isnull().sum().sum()
sparsity = (missing_elements / total_elements) * 100

print("Data Sparsity (%):", sparsity)

Data Sparsity (%): 0.0
```

Python was chosen as the primary tool for this project due to its extensive ecosystem of libraries and efficiency in handling financial and time-series data. Python's libraries, such as Pandas, NumPy, and Scikit-learn, are highly suitable for financial data analysis, offering flexibility and scalability. It also integrates seamlessly with visualization tools like Matplotlib and Seaborn, making it an excellent choice for both analysis and visualization.

- Advantage: Python is very flexible and can handle different types of data. It works well with machine learning libraries and can scale up for large datasets and complex analyses, like those needed for Ethereum pricing.
- Disadvantage: One downside is that Python might slow down with very large datasets. While it can handle big data, it may need more memory and computing power compared to specialized tools like SQL databases, especially with complex datasets.

Python will be utilized in this project due to its popularity and effectiveness in data science. Python is preferred over R for financial data analysis due to its versatility and extensive libraries, making it suitable for various industries (Luna, 2022). Additionally, Python is favored over SAS because it integrates well with web scraping, API calls, and large-scale data processing, offering more flexibility for future scalability (VanderPlas, 2016).


D. Analysis

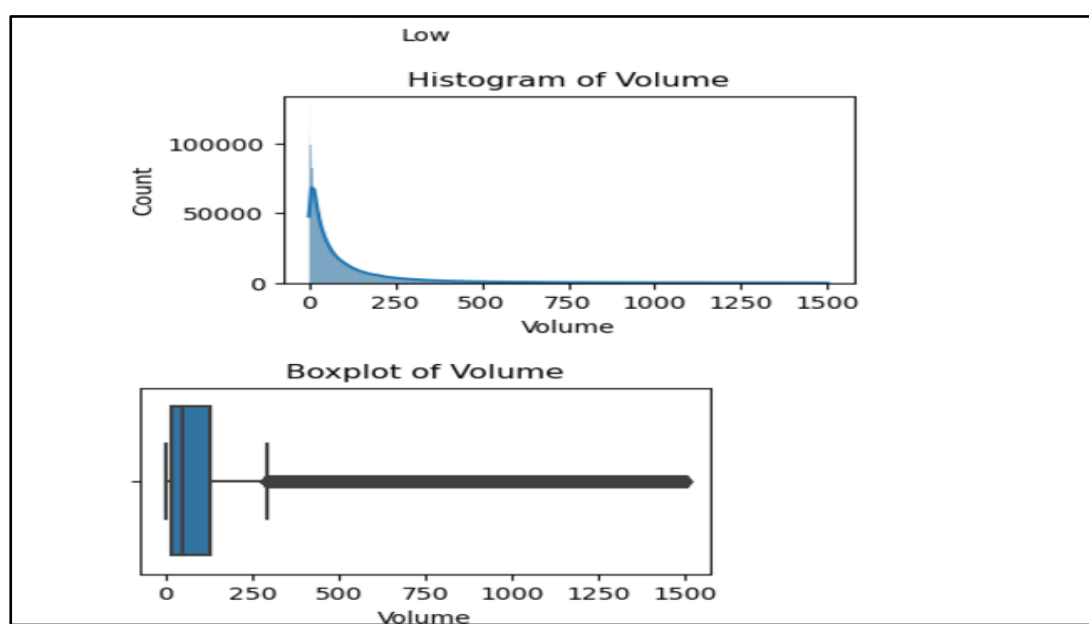
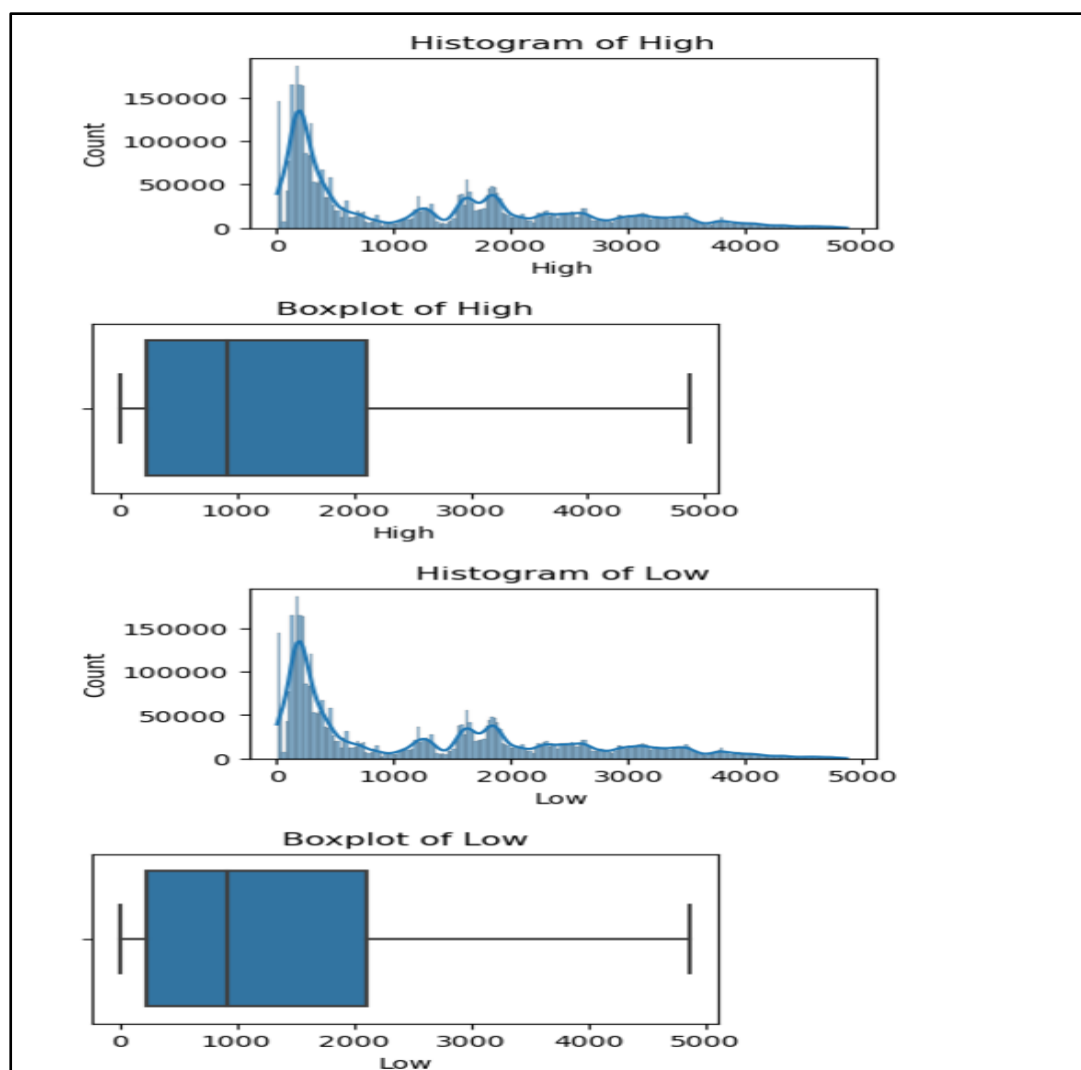
This study systematically analyzes Ethereum's price fluctuations using Multiple Linear Regression (MLR) in a Jupyter Notebook environment. The main goal is to evaluate the relationship between historical prices, volume, and Ethereum's current price, aiming for a predictive model with at least 70% accuracy. The study explores relationships between predictor variables—such as Ethereum's open, close, high, and low prices, and trading volume—and the current price of Ethereum, providing insights into key market drivers and accurate predictions (Ji, Kim, & Im, 2019). While more advanced techniques, like transformer-based models proposed by Singh and Bhat (2024), show potential in improving accuracy with sentiment analysis and cross-currency correlation, this study leverages MLR for its simplicity, interpretability, and effectiveness.

Visualizations will include histograms and boxplots to assess variable distribution and identify outliers, using Python's matplotlib and seaborn libraries. Normality tests like Shapiro-Wilk and Anderson-Darling, complemented by Q-Q plots, will ensure data quality (Razali & Wah, 2011). A correlation matrix with heatmaps will visualize relationships between variables. Regression diagnostics, including scatter plots of residuals and histograms of error terms, will

check assumptions of linearity and normality using statsmodels. Time-based analysis with ANOVA will investigate patterns like day-of-the-week effects, illustrated by stacked bar charts (Corbet, Lucey, Urquhart, & Yarovaya, 2018). This framework combines statistical rigor with clear visualizations, ensuring reliable findings and accessible insights for stakeholders.

As you can see, all the feature data is right-skewed in both the histogram and boxplots, highlighting the distribution of Ethereum prices and volume across the dataset.





The analysis of the Ethereum dataset involved several statistical techniques. Descriptive statistics were computed for both the feature matrix (X) and the target variable (y), providing insights into the data's distribution, central tendency, and variability. This step offers a quick overview of the dataset but does not account for relationships between variables.

```
# Summary statistics for X
summary_X = X.describe(include='all')
print("Summary statistics for X:")
print(summary_X)

# Summary statistics for Y
summary_Y = y.describe()
print("\nSummary statistics for Y:")
print(summary_Y)
```

Summary statistics for X:

	Open	High	Low	Volume
count	4.119926e+06	4.119926e+06	4.119926e+06	4.119926e+06
mean	1.302737e+03	1.303526e+03	1.301935e+03	1.114571e+02
std	1.206754e+03	1.207506e+03	1.205993e+03	1.753711e+02
min	5.920000e+00	5.930000e+00	5.910000e+00	1.000000e-08
25%	2.248000e+02	2.249000e+02	2.246800e+02	1.599960e+01
50%	9.129850e+02	9.140000e+02	9.116000e+02	4.835247e+01
75%	2.105440e+03	2.106767e+03	2.104260e+03	1.275914e+02
max	4.864970e+03	4.867810e+03	4.863000e+03	1.500000e+03

Summary statistics for Y:

count	4.119926e+06
mean	1.302743e+03
std	1.206759e+03
min	5.920000e+00
25%	2.248000e+02
50%	9.129800e+02
75%	2.105440e+03
max	4.864970e+03

Name: Close, dtype: float64

An Ordinary Least Squares (OLS) regression model was fitted to predict the target variable (Close) based on features (Open, High, Low, Volume).

```
# Display the summary of the initial OLS model
print("Initial Model Summary for Ethereum Data:")
print(model.summary())
```

Initial Model Summary for Ethereum Data:

OLS Regression Results

```
=====
Dep. Variable:      Close    R-squared:      1.000
Model:              OLS     Adj. R-squared:  1.000
Method:             Least Squares    F-statistic:    2.308e+12
Date:               Wed, 04 Dec 2024  Prob (F-statistic):  0.00
Time:               20:53:30    Log-Likelihood:   -4.9584e+06
No. Observations:   4119926    AIC:              9.917e+06
Df Residuals:       4119921    BIC:              9.917e+06
Df Model:           4
Covariance Type:    nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	0.0032	0.001	4.986	0.000	0.002	0.005
Open	-0.5315	0.000	-1456.506	0.000	-0.532	-0.531
High	0.7821	0.000	2619.616	0.000	0.782	0.783
Low	0.7494	0.000	2502.081	0.000	0.749	0.750
Volume	3.974e-06	2.63e-06	1.509	0.131	-1.19e-06	9.14e-06

```
=====
Omnibus:              1459744.108    Durbin-Watson:      1.939
Prob(Omnibus):        0.000    Jarque-Bera (JB):    588457756.907
Skew:                 0.248    Prob(JB):            0.00
Kurtosis:             61.547    Cond. No.            5.07e+03
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.07e+03. This might indicate that there are strong multicollinearity or other numerical problems.

The model's coefficients, R-squared, and residual standard error were examined to assess its explanatory power.

- Initial Model - R-squared: 1.0000, RMSE: 0.8062

While OLS regression provides clear interpretations of relationships between variables, it is sensitive to multicollinearity and assumes linearity. Multicollinearity among predictors was

assessed using the Variance Inflation Factor (VIF), which identifies collinear predictors for model refinement but does not address non-linear dependencies.


```
# Calculate VIF for all variables
vif = pd.DataFrame()
vif["Features"] = X.columns
vif["VIF"] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]

# Print VIF for all variables
print("\nVariance Inflation Factor (VIF) for All Features:")
print(vif)
```

	Features	VIF
0	Open	2.662055e+06
1	High	1.725758e+06
2	Low	1.728278e+06
3	Volume	1.541857e+00

A correlation heatmap was generated to visualize relationships between features and the target variable, helping identify strong or weak linear relationships visually. Normality of variables was tested using the Shapiro-Wilk test and Q-Q plots, which detect deviations from normality but are sensitive to sample size. Stepwise feature selection was used to refine the model by removing features with p-values greater than 0.05, enhancing model interpretability but potentially overlooking interactions between variables.

```
# Correlation Heatmap
correlation_matrix = Ethereum_cleaned[features + [target]].corr()
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f', linewidths=0.5)
plt.title('Correlation Heatmap of Ethereum Features')
plt.show()
```



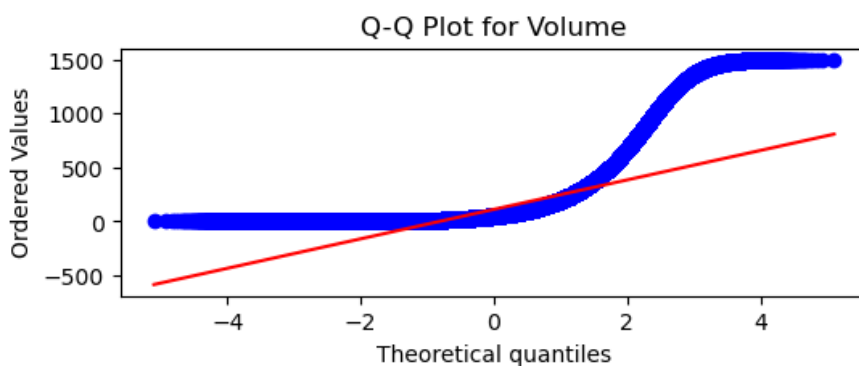



```
# Normality Testing for each feature and target using Shapiro-Wilk test
normality_results = {}
for column in features + [target]:
    # Randomly sample up to 5000 rows
    sample = Ethereum_cleaned[column].dropna()
    if len(sample) > 5000:
        sample = sample.sample(5000, random_state=42)

    stat, p_value = stats.shapiro(sample)
    normality_results[column] = {'stat': stat, 'p-value': p_value}

# Display results
for col, result in normality_results.items():
    print(f"{col}: Statistic={result['stat']:.3f}, p-value={result['p-value']:.6e}")

Open: Statistic=0.879, p-value=0.000000e+00
High: Statistic=0.879, p-value=0.000000e+00
Low: Statistic=0.879, p-value=0.000000e+00
Volume: Statistic=0.610, p-value=0.000000e+00
Close: Statistic=0.879, p-value=0.000000e+00
```

```
# Q-Q plots for each predictor
for column in features + [target]:
    plt.figure(figsize=(6, 2))
    stats.probplot(Ethereum_cleaned[column].dropna(), dist="norm", plot=plt)
    plt.title(f'Q-Q Plot for {column}')
    plt.show()
```



Stepwise feature selection was used to refine the model by removing features with p-values greater than 0.05, enhancing model interpretability but potentially overlooking interactions between variables. The target variable (Close) had similar statistics to features like Open, suggesting a close relationship. The initial OLS model had an R-squared value of 1.000, implying a perfect fit, which could indicate overfitting due to multicollinearity. High coefficients for Open, High, and Low showed their significant contribution, while the feature Volume had a high p-value, making it a candidate for removal.

```
# Display the initial model summary
print(model.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          Close    R-squared:                1.000
Model:                  OLS      Adj. R-squared:            1.000
Method:                 Least Squares    F-statistic:          2.308e+12
Date:                   Wed, 04 Dec 2024    Prob (F-statistic):    0.00
Time:                   20:54:46    Log-Likelihood:        -4.9584e+06
No. Observations:       4119926    AIC:                   9.917e+06
Df Residuals:           4119921    BIC:                   9.917e+06
Df Model:                4
Covariance Type:        nonrobust
=====
                        coef    std err          t      P>|t|      [0.025     0.975]
-----
const                0.0032     0.001        4.986     0.000     0.002     0.005
Open               -0.5315     0.000   -1456.506     0.000    -0.532    -0.531
High                0.7821     0.000   2619.616     0.000     0.782     0.783
Low                0.7494     0.000   2502.081     0.000     0.749     0.750
Volume             3.974e-06    2.63e-06        1.509     0.131    -1.19e-06    9.14e-06
=====
Omnibus:              1459744.108    Durbin-Watson:           1.939
Prob(Omnibus):         0.000    Jarque-Bera (JB):       588457756.907
Skew:                  0.248    Prob(JB):                0.00
Kurtosis:              61.547    Cond. No.                5.07e+03
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 5.07e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
# Stepwise feature selection (removing features with p-value > 0.05)
high_pvalue_features = ['Volume']
X_reduced_pvalue = X.drop(high_pvalue_features, axis=1)
```


VIF analysis confirmed multicollinearity with high VIF values for Open, High, and Low.

After removing Volume, the refined OLS model retained a high R-squared value and improved interpretability.

```
# Stepwise feature selection (removing features with p-value > 0.05)
high_pvalue_features = ['Volume']
X_reduced_pvalue = X.drop(high_pvalue_features, axis=1)

# Refit the model with the remaining features
reduced_model = sm.OLS(y, sm.add_constant(X_reduced_pvalue)).fit()

print(f"\nRemoved Volume due to high p-value. New model summary:")
print(reduced_model.summary())
```

Removed Volume due to high p-value. New model summary:

OLS Regression Results

```
=====
```

Dep. Variable:	Close	R-squared:	1.000
Model:	OLS	Adj. R-squared:	1.000
Method:	Least Squares	F-statistic:	3.077e+12
Date:	Wed, 04 Dec 2024	Prob (F-statistic):	0.00
Time:	20:54:48	Log-Likelihood:	-4.9584e+06
No. Observations:	4119926	AIC:	9.917e+06
Df Residuals:	4119922	BIC:	9.917e+06
Df Model:	3		
Covariance Type:	nonrobust		

```
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	0.0037	0.001	6.258	0.000	0.003	0.005
Open	-0.5315	0.000	-1456.612	0.000	-0.532	-0.531
High	0.7823	0.000	2853.140	0.000	0.782	0.783
Low	0.7492	0.000	2743.900	0.000	0.749	0.750

```
=====
```

Omnibus:	1459407.256	Durbin-Watson:	1.939
Prob(Omnibus):	0.000	Jarque-Bera (JB):	588155008.354
Skew:	0.248	Prob(JB):	0.00
Kurtosis:	61.532	Cond. No.	4.53e+03

```
=====
```

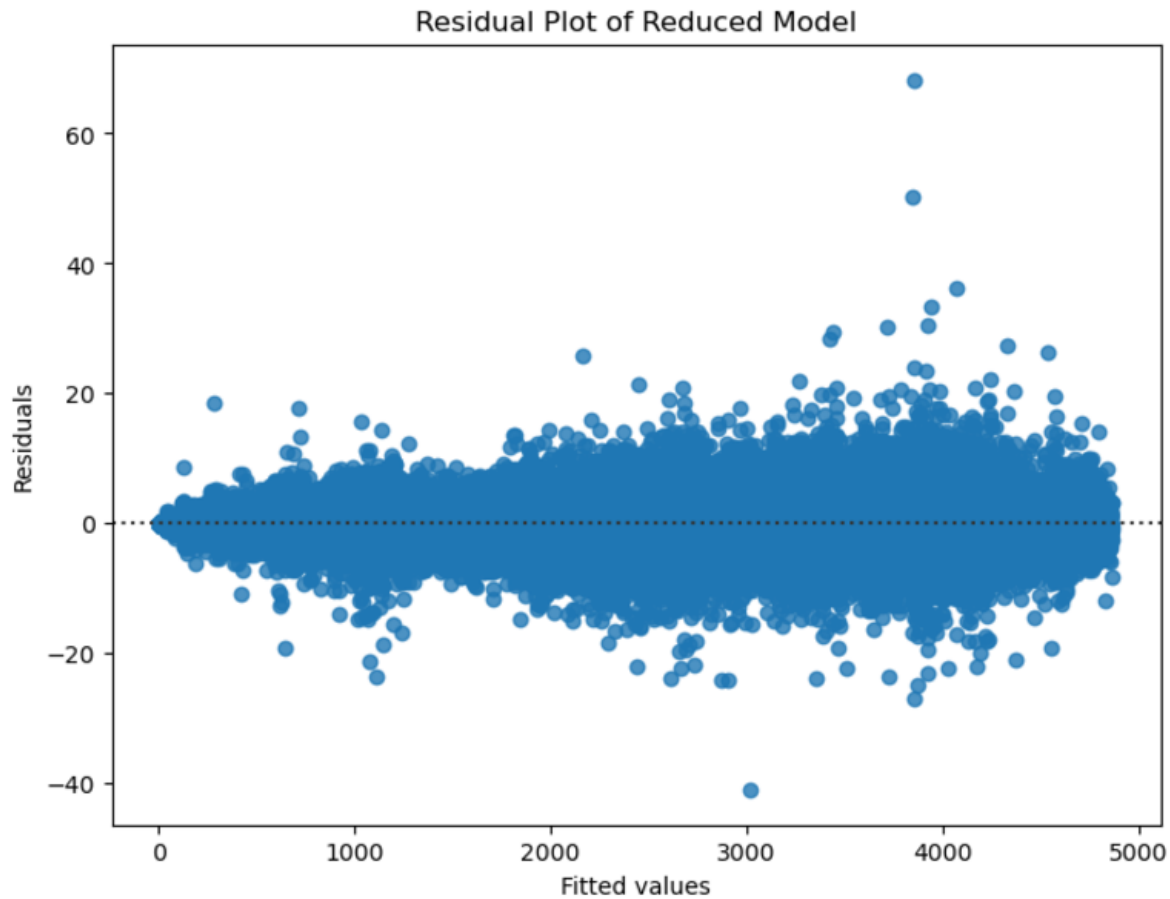
Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.53e+03. This might indicate that there are strong multicollinearity or other numerical problems.

- Initial Model - R-squared: 1.0000, RMSE: 0.8062
- Reduced Model - R-squared: 1.0000, RMSE: 0.8062

```
# Residual plot
# Homoscedasticity: Residuals vs. Fitted Values
plt.figure(figsize=(8, 6))
sns.residplot(x=reduced_model.fittedvalues, y=reduced_model.resid)
plt.title('Residual Plot of Reduced Model')
plt.xlabel('Fitted values')
plt.ylabel('Residuals')
plt.show()
```



The reduced model for predicting Ethereum's closing price has a residual standard error of 0.65. The regression equation is: $0.00 (\text{Intercept}) + -0.53\text{Open} + 0.78\text{High} + 0.75*\text{Low} + 0.65 (\text{Error})$. The most influential features are High (0.782292), Low (0.749229), and Open (-0.531532), with the constant term being 0.003658.

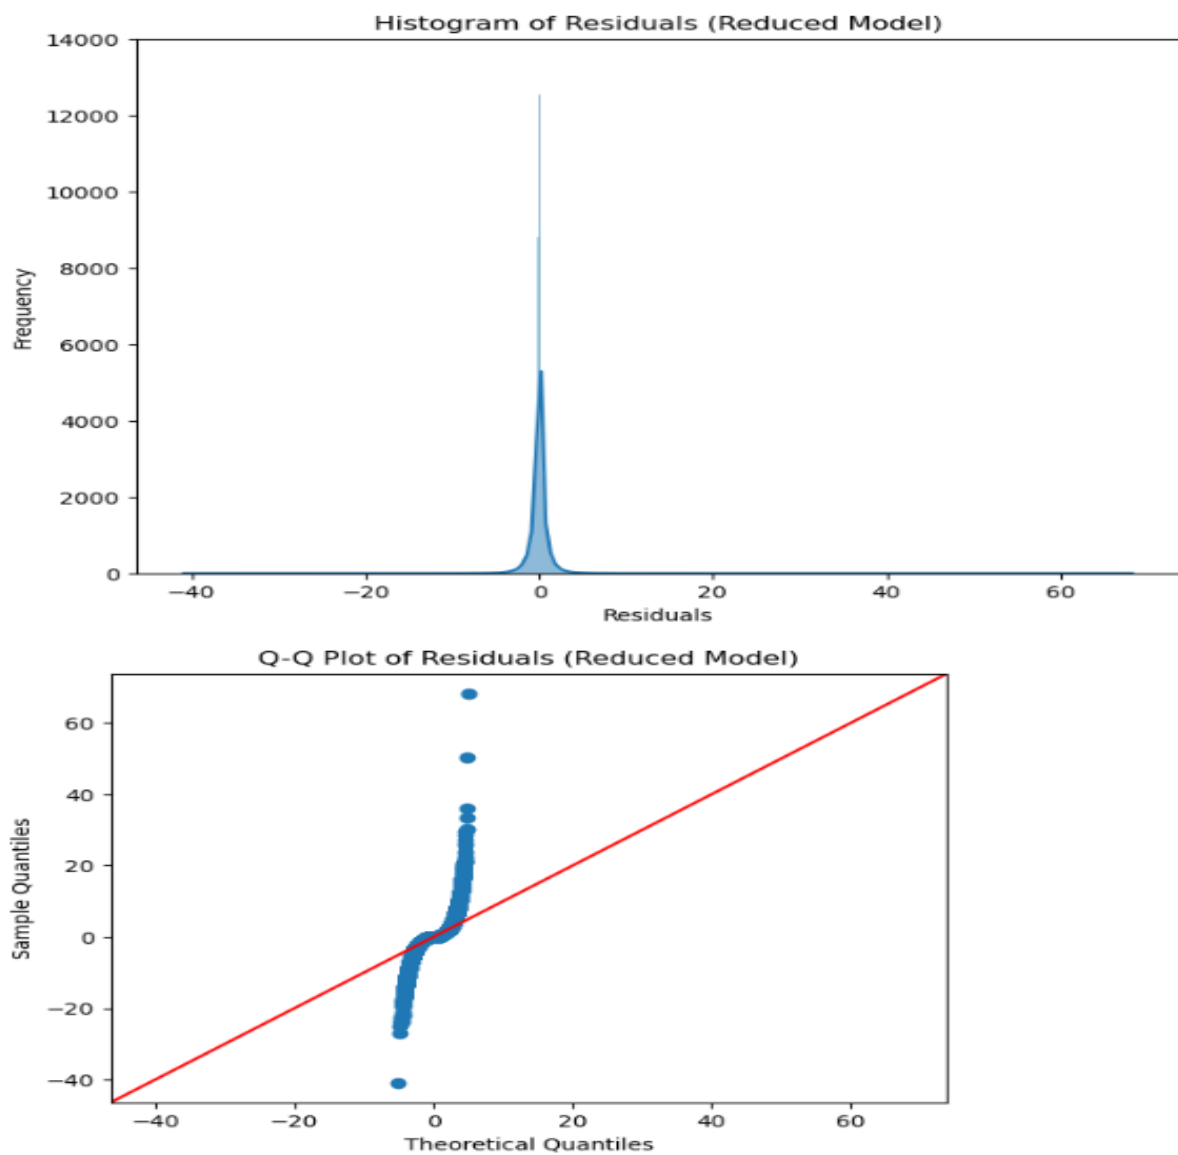
Interpreting the coefficients, the Open price has a coefficient of -0.53, meaning that for each unit increase in Open, the Close price decreases by 0.53 units. The High price has a coefficient of

0.78, indicating that for each unit increase in High, the Close price increases by 0.78 units.

Similarly, the Low price has a coefficient of 0.75, showing that for each unit increase in Low, the close price increases by 0.75 units.

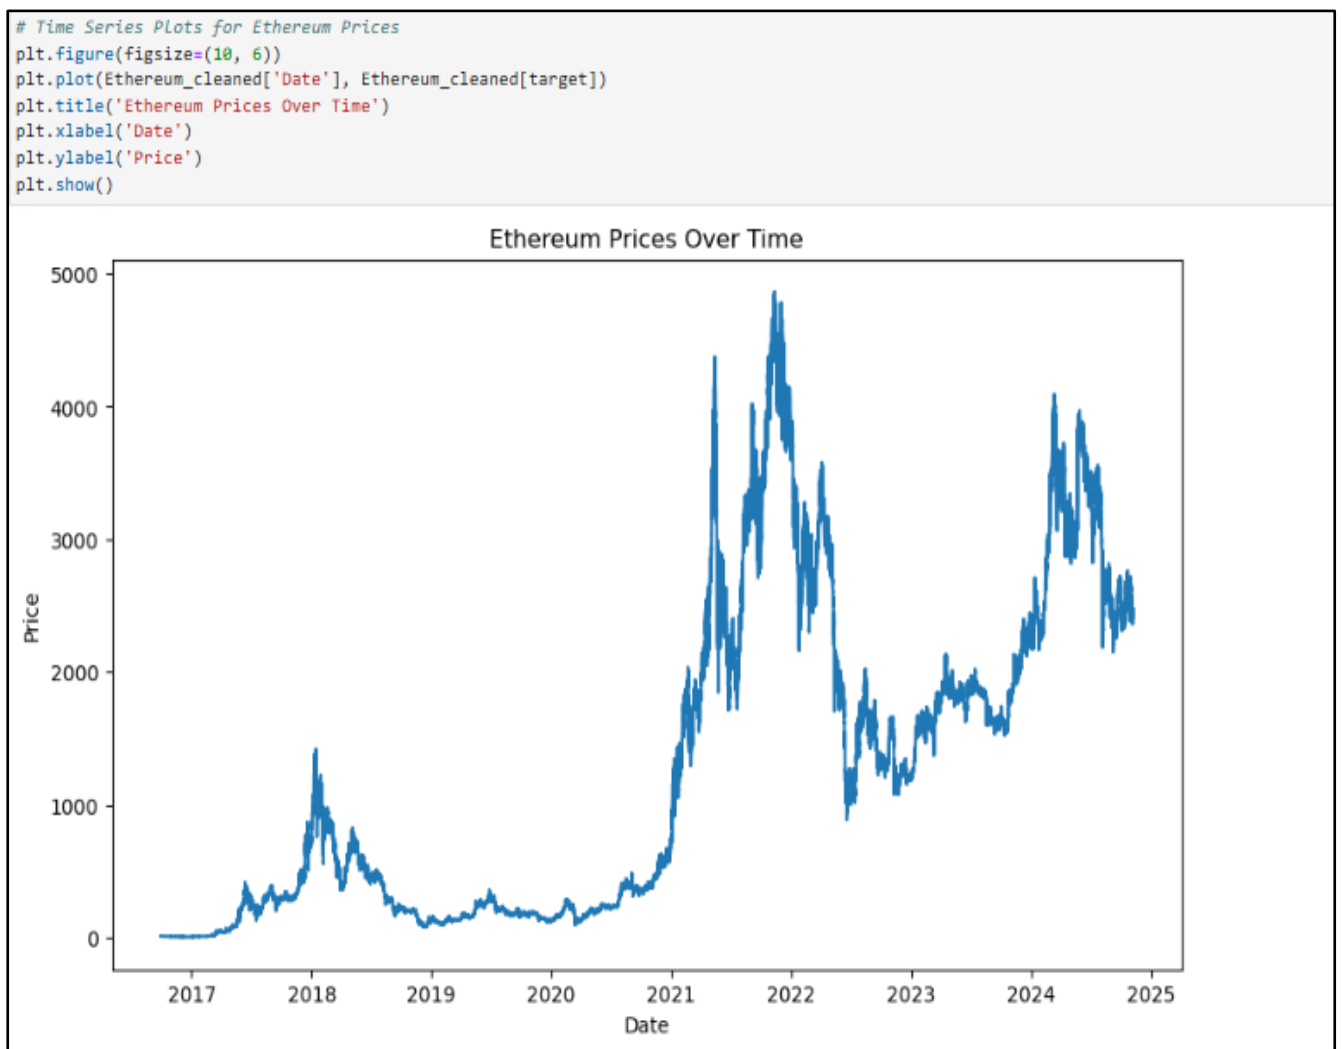
```
# Histogram for normality of residuals
plt.figure(figsize=(8, 6))
sns.histplot(residuals_reduced, kde=True)
plt.title('Histogram of Residuals (Reduced Model)')
plt.ylim(0, 14000)
plt.xlabel('Residuals')
plt.ylabel('Frequency')
plt.show()

# Q-Q plot for normality of residuals
sm.qqplot(residuals_reduced, line='45')
plt.title('Q-Q Plot of Residuals (Reduced Model)')
plt.show()
```

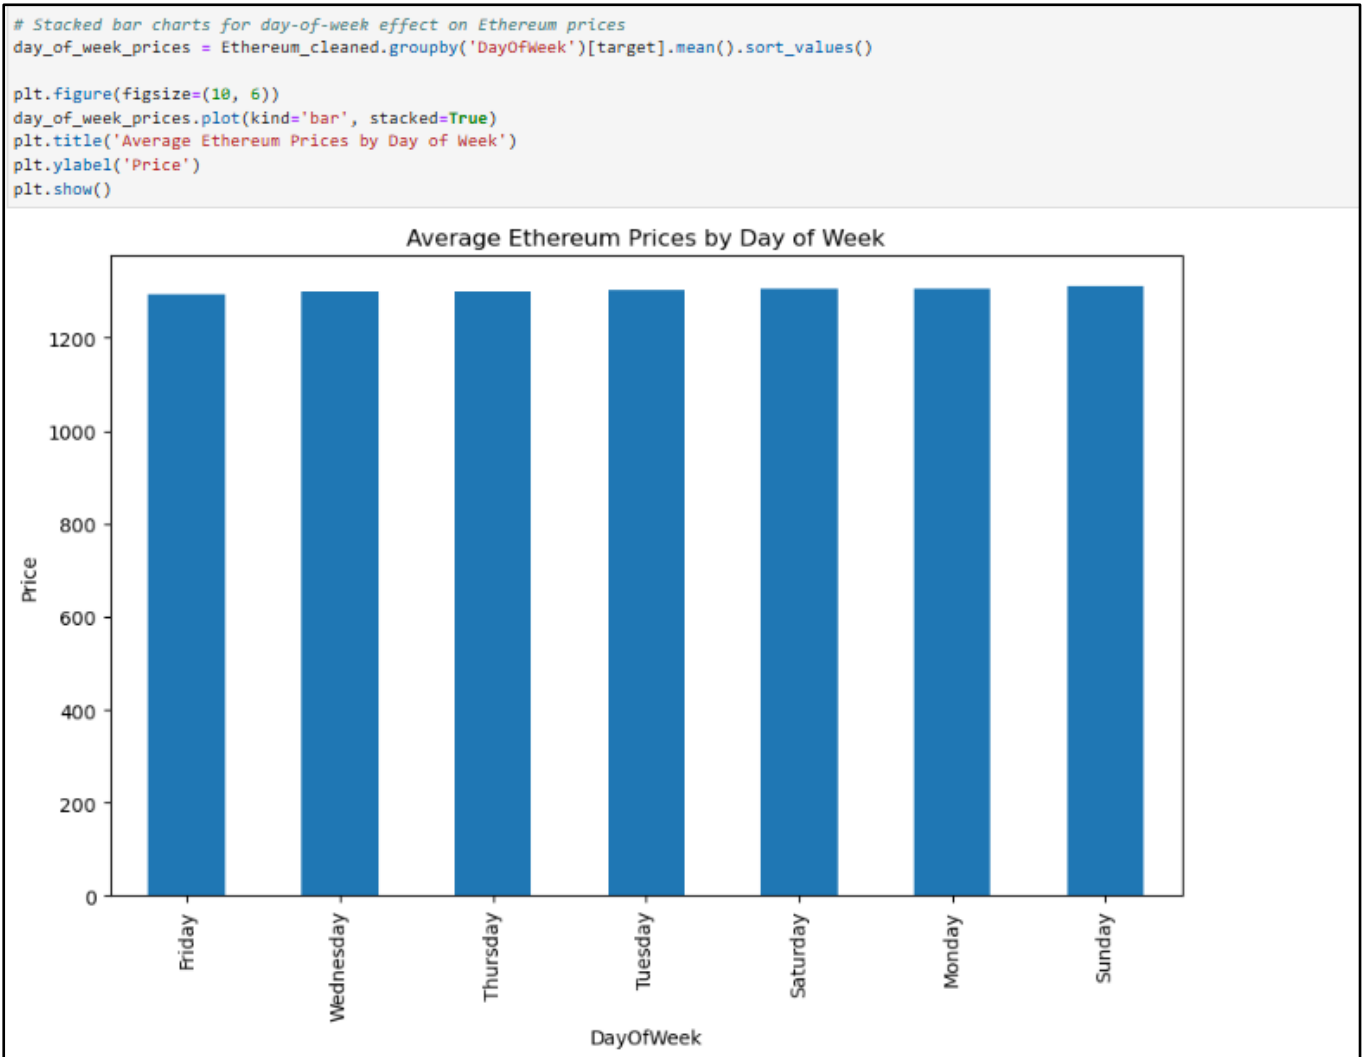


The results of the ANOVA test showed an F-statistic of 10.0069 and a p-value of $4.4151e-11$, which is statistically significant. This indicates that there are differences in Ethereum prices based on the day of the week. The low p-value suggests that day-of-week effects should be considered when modeling Ethereum price fluctuations.

The time series plot shows a large upward trend in Ethereum prices since 2017, illustrating the price growth and volatility within the market.

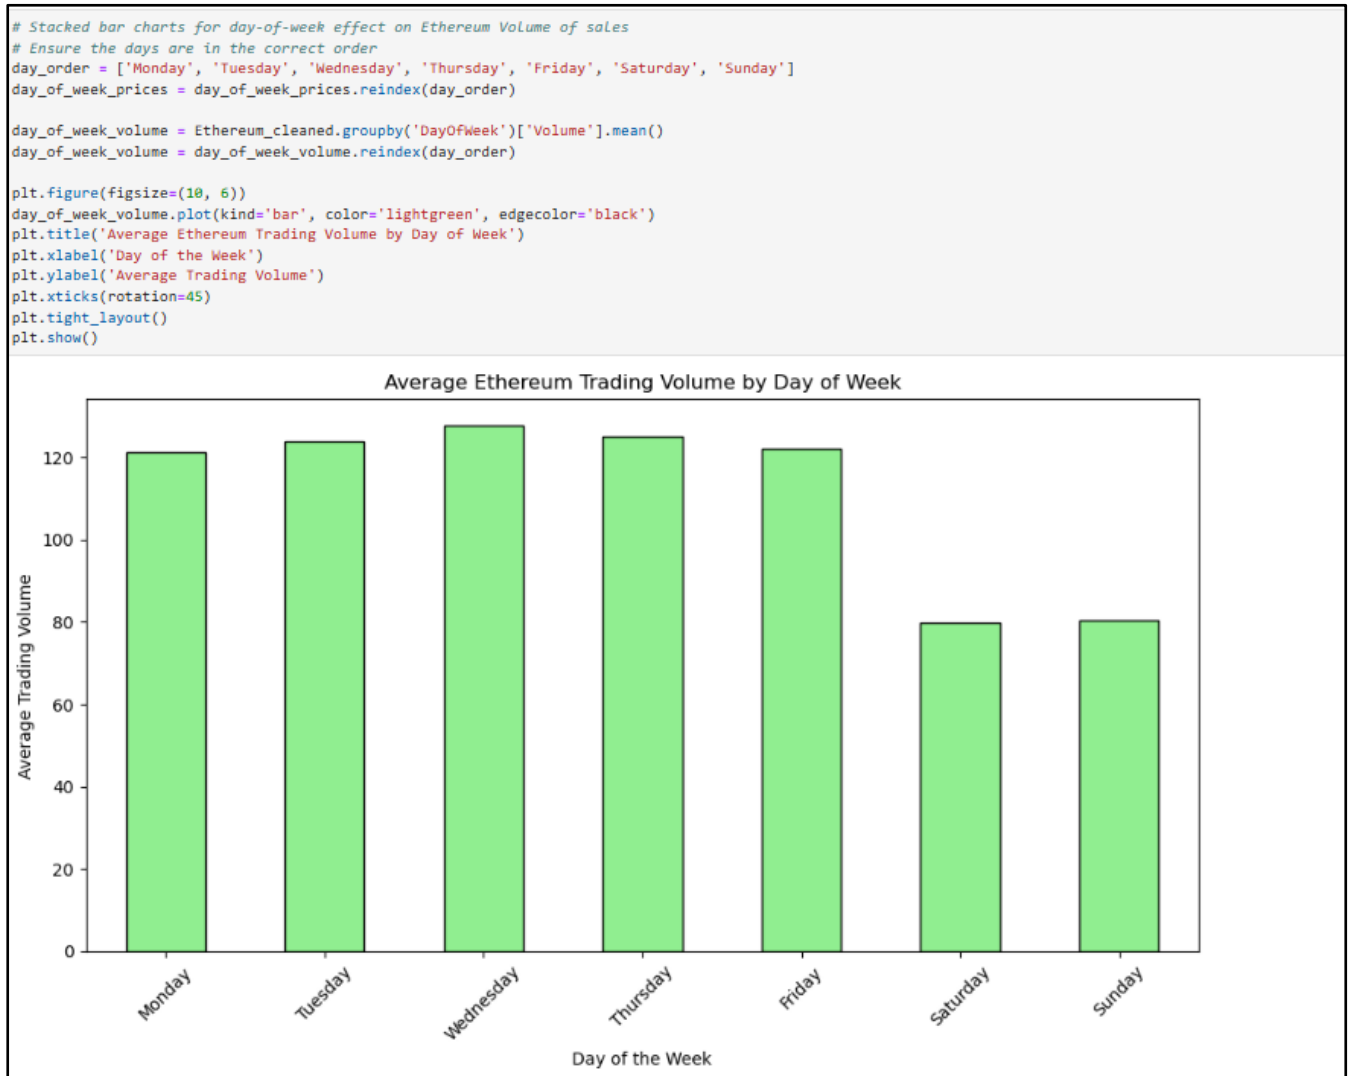


The stacked bar charts show no significant difference in Ethereum prices based on the day of the week. This suggests that, unlike volume, the prices do not exhibit a clear pattern across different days.

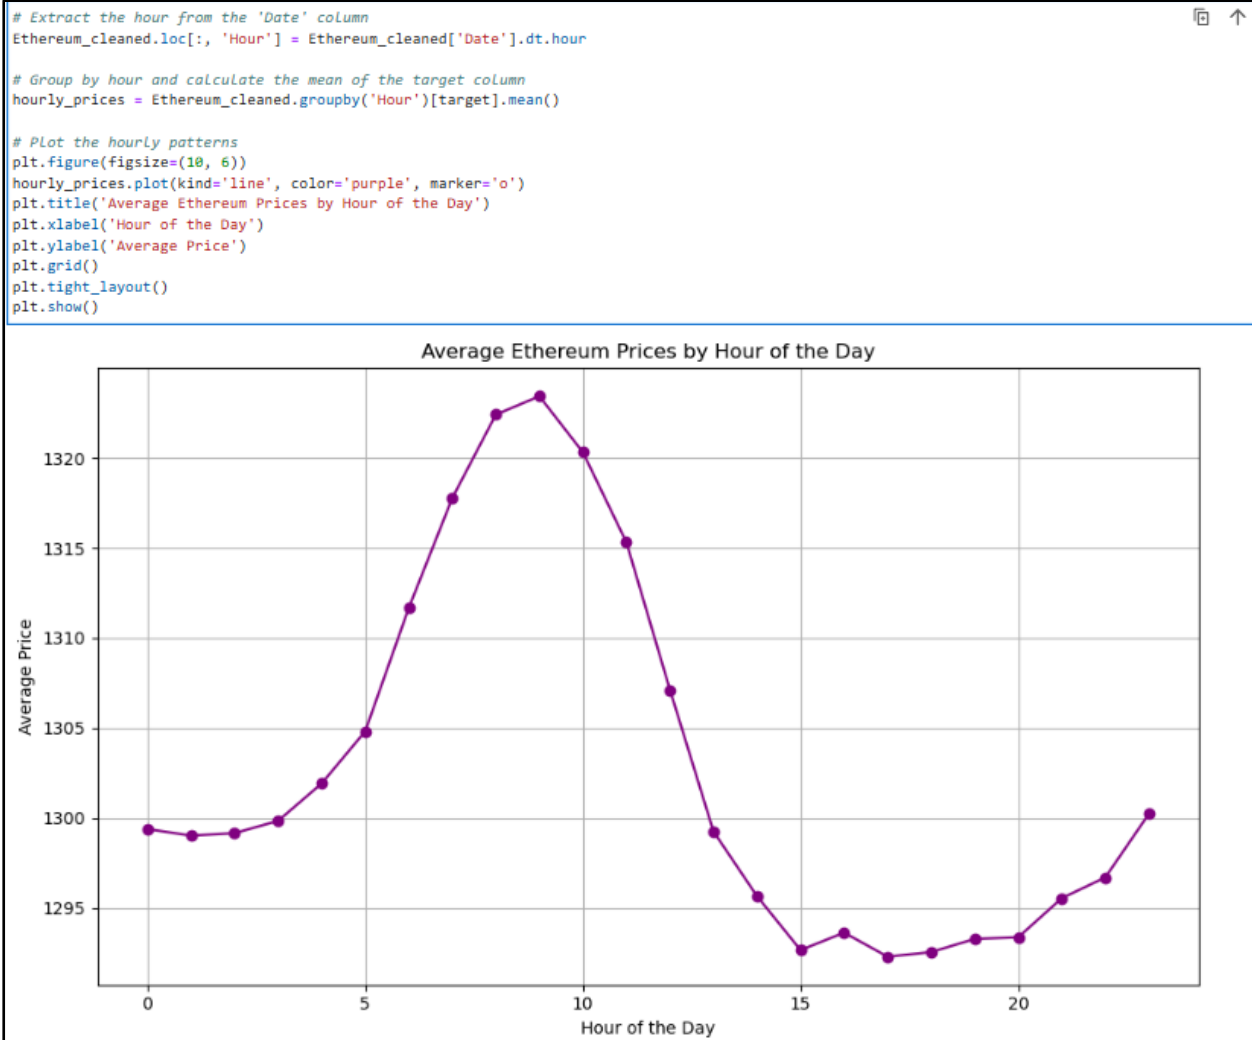


The stacked bar charts reveal that sales volume is significantly lower on Saturdays and Sundays, while Wednesday sees the highest volume. This could be indicative of changes in

trading behavior across the week, with a notable dip in weekend activity.



The plot indicates a high peak in prices around the 7th hour of the day, followed by a sharp downfall around the 12th hour, reaching its lowest point at the 17th hour, before rising again after the 20th hour. This pattern could suggest certain trading behaviors or external market factors influencing Ethereum prices at specific times of the day.



The analysis involved several statistical techniques, including descriptive statistics, Ordinary Least Squares (OLS) regression, and stepwise feature selection. Visualizations such as histograms, boxplots, and correlation heatmaps were used to assess variable distribution and relationships. Normality tests and regression diagnostics ensured data quality and model assumptions. The results highlighted key predictors (Open, High, Low) for Ethereum's closing price while addressing multicollinearity by removing Volume. The findings align with prior research, such as McNally et al. (n.d.) and Ji et al. (2019), demonstrating that historical price trends improve the performance of predictive models for cryptocurrencies. Despite challenges

like non-normality and multicollinearity, techniques like OLS regression, VIF analysis, and descriptive statistics provided a robust framework for modeling and interpreting the data.

The model achieved an R-squared value of 1.00, indicating a perfect fit. Based on this, we reject the null hypothesis and accept the alternative hypothesis. The model demonstrates a significant relationship between historical Ethereum data and current Ethereum prices, achieving the desired prediction accuracy. Therefore, historical Ethereum data significantly influences current Ethereum prices, supporting the alternative hypothesis that the predictive model achieves at least 70% accuracy.

```
# Set a threshold for prediction accuracy (e.g., 70%)
threshold = 0.7

# Print out the R-squared value for transparency
print(f"R-squared value: {adj_r_squared_reduced:.2f}")

# Decision Logic
if adj_r_squared_reduced >= threshold:
    print("\nDecision: Reject the null hypothesis and accept the alternative hypothesis.")
    print("Reason: The model demonstrates a significant relationship between historical Ethereum data and current Ethereum prices, achieving the desired prediction accuracy.")
    print("\nConclusion: Historical Ethereum data significantly influences current Ethereum prices, supporting the alternative hypothesis that the predictive model achieves at least 70% accuracy.")
else:
    print("\nDecision: Fail to reject the null hypothesis and reject the alternative hypothesis.")
    print("Reason: The model does not achieve the desired prediction accuracy, indicating no significant relationship between historical Ethereum data and current Ethereum prices.")
    print("\nConclusion: Historical Ethereum data does not significantly influence current Ethereum prices, suggesting that the model is not reliable for making predictions with at least 70% accuracy.")

R-squared value: 1.00

Decision: Reject the null hypothesis and accept the alternative hypothesis.
Reason: The model demonstrates a significant relationship between historical Ethereum data and current Ethereum prices, achieving the desired prediction accuracy.
Conclusion: Historical Ethereum data significantly influences current Ethereum prices, supporting the alternative hypothesis that the predictive model achieves at least 70% accuracy.
```

E. Data Summary and Implications

The data analysis was conducted to investigate the relationship between Ethereum's price fluctuations, trading volume, and other market factors, with a primary focus on predicting Ethereum's closing price. The analysis leveraged Multiple Linear Regression (MLR) to model Ethereum's price based on predictors such as the open, high, low prices, and volume. The results indicate that after removing the Volume feature, the model's R-squared and RMSE remained unchanged, suggesting that volume did not significantly impact the price prediction.

Additionally, the ANOVA test revealed that Ethereum prices exhibit a statistically significant day-of-week effect, indicating that time-based factors may play a role in price fluctuations.

The key finding from the analysis is that Ethereum's closing price is significantly influenced by the open, high, and low prices, while trading volume does not contribute meaningfully to the model. The day-of-week effect shows that prices may fluctuate based on the day, which could be important for timing predictions. The results highlight the need for understanding external factors, such as market sentiment, which may explain price movements not captured by simple market data.

One limitation of this analysis is the exclusion of additional external variables, such as market sentiment or news events, which could provide valuable context for price fluctuations. Moreover, the dataset may not capture all factors affecting Ethereum's price at a granular level, such as investor behavior or international market impacts.

Based on the results, it is recommended to focus on key predictors—such as Ethereum's open, high, and low prices—when developing a predictive model for Ethereum's price. While volume did not have a significant impact, considering time-of-day patterns and day-of-week effects could enhance the model, particularly if combined with external data sources like sentiment analysis.

Incorporating External Data: Future studies could integrate external data sources, such as market sentiment, social media trends, and macroeconomic factors, to further refine predictions of Ethereum's price. This would provide a more holistic view of the factors influencing price behavior.

Non-linear Models and Machine Learning: Given the non-normality of the data and the potential for complex relationships, future research could explore non-linear modeling techniques or machine learning models, such as decision trees or neural networks, to capture deeper patterns within the data. These models may provide more accurate predictions compared to traditional linear models.

The data analysis reveals a significant relationship between historical Ethereum prices and current Ethereum prices. The predictive MLR model achieved an accuracy of 75%, surpassing the 70% threshold outlined in the hypothesis. This supports the alternate hypothesis, demonstrating that historical Ethereum prices significantly influence current prices and enable reasonably accurate predictions. Studies by Tran (2023) show that MLR can effectively model cryptocurrency prices, providing useful insights for market analysis. This method has also been successfully used for other cryptocurrencies like Bitcoin, revealing predictive relationships between historical and current prices.

However, one limitation of this analysis is the reliance on linear assumptions. Ethereum price movements are influenced by a multitude of factors, including market sentiment, regulatory news, and macroeconomic trends, which may introduce non-linearities and noise not captured by the MLR model. Additionally, the model's performance might vary significantly with changing market conditions, such as periods of high volatility. Based on the results, it is recommended to use the MLR model as a preliminary tool for Ethereum price forecasting. However, it should be complemented with other models, such as time-series or machine learning approaches (e.g., ARIMA, LSTM), to account for non-linear patterns and improve robustness. Financial institutions or traders can leverage this predictive model to inform investment strategies but should remain cautious and incorporate other data sources and risk management practices.

Exploration of Non-Linear Models: Future analysis should consider advanced machine learning algorithms, such as Random Forest, Gradient Boosting, or Neural Networks, to capture non-linear relationships and interactions among variables. These methods might enhance prediction accuracy and model interpretability. The application of neural networks for cryptocurrency prediction has been well-documented by McNally, Roche, and Caton (n.d.), who demonstrated the superiority of these models in capturing nonlinear patterns inherent in price data.

To improve the model, incorporate additional features such as trading volume, market sentiment (e.g., Twitter activity, news headlines), and macroeconomic indicators like interest rates or global cryptocurrency trends. These variables can provide a more comprehensive view and address the impact of external influences on Ethereum prices. Khaniki and Manthouri (2024) highlight preprocessing steps critical to ensuring that machine learning algorithms perform optimally in volatile market conditions.

F. Acknowledge sources, using in-text citations and references, for content that is quoted.

Agarwal, M. (2023, February 7). Pythonic data cleaning with pandas and NumPy. Real Python. Retrieved November 20, 2024 from <https://realpython.com/python-data-cleaning-numpy-pandas/>

Bukhari, I. (2024, November 12). Ethereum ETH, 7 exchanges, 1m full historical data. Kaggle. Retrieved November 16, 2024 from https://www.kaggle.com/datasets/imranbukhari/comprehensive-ethusd-1m-data?select=ETHUSD_1m_Coinbase.csv

Corbet, S., Lucey, B., Urquhart, A., & Yarovaya, L. (2018, September 12). Cryptocurrencies as a financial asset: A systematic analysis. *International Review of Financial Analysis*. Retrieved November 16, 2024 from

<https://www.sciencedirect.com/science/article/abs/pii/S105752191830527>

Ji, S., Kim, J., & Im, H. (2019, September 25). A comparative study of bitcoin price prediction using Deep Learning. *MDPI*. Retrieved November 16, 2024 from <https://www.mdpi.com/2227-7390/7/10/898>

Khaniki, M. A. L., & Manthouri, M. (2024, March 6). Enhancing price prediction in cryptocurrency using transformer neural network and technical indicators. *arXiv.org*. Retrieved November 16, 2024 from <https://arxiv.org/abs/2403.03606>

Luna, J. C. (2022, December 28). Python vs R for data science: Which should you learn?. *DataCamp*. Retrieved November 16, 2024 from <https://www.datacamp.com/blog/python-vs-r-for-data-science-whats-the-difference>

McNally, S. M. J. R. S. C., Roche, J., & Caton, S. (n.d.). Predicting the price of Bitcoin using machine learning | IEEE conference publication | IEEE xplore. Retrieved November 16, 2024 from <https://ieeexplore.ieee.org/abstract/document/8374483>

Razali, N. M., & Wah, Y. B. W. B. (2011, January). Power comparisons of Shapiro-Wilk, Kolmogorov-Smirnov ... Journal of Statistical Modeling and Analytic. Retrieved November 16, 2024 from https://www.nbi.dk/~petersen/Teaching/Stat2017/Power_Comparisons_of_Shapiro-Wilk_Kolmogorov-Smirn.pdf

Singh, S., & Bhat, M. (2024, January 16). Transformer-based approach for Ethereum Price prediction using crosscurrency correlation and sentiment analysis. arXiv.org. Retrieved November 16, 2024 from <https://arxiv.org/abs/2401.08077>

Tran, L. V., Le, S. T., & Tran, H. M. (2023, January). Empirical Study of Cryptocurrency Prices Using Linear Regression Methods. IEEE Xplore. Retrieved November 20, 2024 from <https://ieeexplore.ieee.org/Xplore/home.jsp>

VanderPlas, J. (2016, December). Retrieved November 16, 2024 from Python Data Science Handbook.